

NITHISH KADAM GANESH

Boston, MA | (617) 415-6651 | nithishkadam@gmail.com | www.nithishkadam.com

Dear Hiring Team,

I am completing an M.S. in Computer Science at Northeastern University (3.9 GPA, December 2026), with two and a half years of prior full-time engineering experience and an AWS Solutions Architect certification.

At Manhattan Associates I designed a message processing framework that orchestrates warehouse fulfillment for dozens of enterprise clients globally. It cut end-to-end latency by 30% at 5,000+ orders per minute and enforced strict ordering through partition keys, so picking always precedes packing and replenishment always precedes the next picking wave. I owned it in production: 150+ incidents resolved across 30+ deployments. That is the work I want to keep doing, systems that have to work under load, under pressure, and across teams.

The hardest thing I built at Manhattan was a barcode parsing and validation framework for a global enterprise client. Field-scanned barcodes arrived in 12 distinct formats, each packing weight, temperature zones, item identifiers, and sequencing into a single code, and every scan triggered 7 downstream transactions inside their WMS. I spent time reverse-engineering the formats from real field data and sitting with the functional consultant team to learn every validation rule, because a parsing error meant product in the wrong temperature zone and spoiled inventory. I was recognized with the company's "Above & Beyond" and "Star of the Quarter" awards.

At Wintergreen, a startup, I introduced an entire healthcare reporting stack from scratch and owned architecture decisions across all four layers: SQL data shaping, Node/Express APIs, JasperReports, and a React Native client. I also designed RESTful APIs to bring legacy PHP/Laravel systems into that pipeline.

I build outside of work, constantly. AlphaWealth is a self-hosted, GPU-accelerated finance AI platform of 10+ microservices across Java, Python, C++, and React, with an advisor that orchestrates four hot-swappable LLM backends over a RAG pipeline.

Arcanum (arcanumverse.org) is a story-driven puzzle/CTF platform a few of us pitched to Northeastern and shipped for a week-long campus-wide event: 140+ registered users, 90+ concurrent, and 12,000+ submissions on day one. I owned three of the ten trials, contributed to the semantic scoring engine that graded answers by meaning rather than exact match, and patched a live leaderboard defect mid-event with zero downtime.

I care about systems that are reliable, observable, and built to scale, whether the problem is backend infrastructure, API design, or full-stack delivery. I would welcome the chance to talk.

Best regards,
Nithish Kadam Ganesh